

PUBLIC RELEASE · MAY 2026

The Citizen Think Tank

How AI Closes the Power Gap in Democratic Policy Production

Reform Democrats of the New Commonwealth (RDNC) · May 2026

Authors

An American Citizen

Independent — no institutional affiliation

Claude (Anthropic)

AI research and drafting system

ABSTRACT

This paper documents and analyzes a proof of concept: the production of a comprehensive, institutionally structured policy framework by a single private citizen with no advanced degree, no government experience, and no institutional affiliation, working with AI as the primary research and drafting tool. The output — 65 sourced and scored policy proposals, 12 bills drafted in legislative language, a three-phase implementation sequence, an economic feasibility analysis across three scenarios, a constitutional architecture analysis across ten policy domains integrated into operative bill text, and an approximately 50-document supporting ecosystem — is examined as evidence of a structural shift in who can produce credible democratic policy.

The paper argues four things in sequence. First, that the form of credible policy production — the specific institutional capacity that organizations like the Heritage Foundation, the American Enterprise Institute, and the Center for American Progress have monopolized — has been functionally replicated by one person with consumer AI tools, at a direct cost of approximately \$150, in working sessions measured in weeks. Second, that what AI specifically replaced is not "expertise" or "time" generically, but the specific institutional capacity to produce, vet, source, sequence, and translate policy at the level of specificity that legislative and political systems take seriously — and that this displacement can be demonstrated rather than asserted. Third, that the project's most contested methodological problem — AI helping to draft AI legislation — was answered in operative form rather than disclaimed away, and that the response provides a template for honest AI-assisted policy work where the AI is not a neutral party. Fourth, that the democratic implication is direct: broad, affordable access to capable AI is now a precondition for democratic participation, and the concentration of that access would amplify the asymmetries it has temporarily reduced.

PART I

The Proof

I.1 The Problem in One Number

The Heritage Foundation operates on an annual budget exceeding \$110 million and employs more than 500 people. Its flagship policy product for the 2025 presidential transition, *Mandate for Leadership*, represents the accumulated work of hundreds of contributors over years. The American Enterprise Institute operates on approximately \$70 million annually. The Brookings Institution, approximately \$100 million. The Center for American Progress, the closest Democratic equivalent in form and function, operates on roughly \$50 million.

These are not extravagant operations by Washington standards. They represent the historical minimum viable institutional footprint for producing policy at the level of specificity that governments, candidates, journalists, and coalition partners take seriously. The number — \$110 million for Heritage — is not just a budget. It is the price of admission to a particular kind of democratic participation: the kind in which one's policy ideas reach the systems where law actually gets made, with sufficient rigor and documentation to survive contact with those systems.

That price has, until very recently, been treated as a fixed feature of democratic political infrastructure. Citizens could vote, organize, protest, and write letters. They could not produce the kind of structured, sourced, legislatively translated policy document that institutional actors produced. The asymmetry was not a bug in the democratic system. It was load-bearing structure.

This paper documents what happens when the load-bearing structure stops bearing load.

I.2 The Origin

The American Renewal Agenda (ARA) began as a single document drafted by a private American citizen before any AI collaboration. It is approximately 3,000 words, organized loosely under category headings — Legal, Infrastructure, Military, Healthcare and Education and Social, Foreign Affairs, Agriculture, and a residual category — and contains roughly 80 distinct policy ideas. It is unedited and unsourced.

That document — preserved in the project archive as the genesis text — is the seed of everything that follows. It is the proof of concept's most important artifact, because it establishes what the input actually was. No prior policy training. No institutional affiliation. No staff. A citizen who had been paying attention, sitting down to write what he thought.

What turned that document into an institutional infrastructure was a sequence of working sessions with an AI system over several weeks. The author did not delegate. He did not paste the document into a model and ask for a finished framework. The collaboration was iterative and substantive: structural decisions about how to organize a four-layer architecture, design choices about how to score proposals against international evidence, judgment calls about what belonged in the framework and what did not, strategic decisions about how the ecosystem would be released. The AI contributed research synthesis, legislative translation, constitutional architecture analysis, and consistency maintenance across what eventually became an approximately 50-document ecosystem. Neither party could have produced the result alone.

The genesis document still exists. The contrast between it and the final ecosystem is the most direct evidence the proof of concept can offer. One person who could write a coherent list of complaints became, over weeks, the author of record on a policy framework with the institutional form that gets taken seriously.

I.3 The Architecture

The American Renewal Agenda is structured as a four-layer document ecosystem, reflecting the different audiences and decision rights that policy production serves.

The reference layer holds the production-grade analytical and legislative work: the Framework v26 with 65 sourced and scored proposals across three implementation phases, the 12-bill series in legislative language, the Scoring Methodology and Scoring Justification Record (195 axis paragraphs), the Source Reference Guide with 370 primary sources, the Capital Report with three-scenario economic feasibility analysis and 141 tier-rated citations, the Constitutional Architecture analysis across ten domains, and the Personnel Pipeline.

The internal layer holds the governance, operational, and strategic infrastructure: the Project Ecosystem Index, the Document Ecosystem Index with per-bill Roster, the Editorial Protocol, the Brand Bridge messaging discipline, the Release Architecture, the Structural Independence Doctrine, the IdeaNotepad, and the AI Working Notes that captures cross-session operational knowledge.

The public RDNC layer holds institutional-voice documents: the Ideological Toolkit, the Framework Paper, and the present white paper. The public ARI layer holds the candidate- and voter-facing campaign infrastructure: the One-Pager, the Policy Platform, the Attack-Answer Companion, the Timed Pitches library, and the Candidate Orientation guide. A public comparison layer holds the ARA-versus-Project-2025, ARA-versus-DNC-2024, and three-way analysis documents.

The architecture is not decoration. It is what allows the work to function institutionally — the same content available at different registers for different audiences, with the consistency of analysis preserved across translations.

I.4 The Output

The institutional comparison places these numbers in scale.

Dimension	ARA (this project)	Project 2025	DNC 2024 Platform
Policy proposals scored	65	37 analyzed	33 analyzed
Average evidence score	7.70	4.01	7.76
Legislative bills drafted	12	0 (mandate-level)	0
Primary sources cited	370	~40 (estimated)	~25 (estimated)
Constitutional architecture analysis	10 domains, integrated into operative bill text	Limited	Limited
Economic feasibility analysis	3 scenarios, 141 tier-rated citations	Not provided	Not provided
Institutional budget	~\$150 out-of-pocket	\$22M+ estimated (Heritage: \$110M annually)	\$8M+ estimated for platform development
Production time	Weeks	4+ years	18+ months
Staff	1 person + AI	500+ contributors	Institutional staff

These figures are not offered as evidence of superiority across all dimensions. Institutional think tanks have access to political networks, proprietary data, legal review cycles, and operational experience that AI-assisted individual production cannot replicate. The comparison establishes scale and structural equivalence — that the form of credible policy production has been replicated, not that it has been exceeded.

The structural equivalence is the proof. The form through which policy ideas become political reality is no longer proprietary to well-funded institutions. Whether it remains accessible to citizens depends on what happens next.

PART II

The Mechanisms

II.1 What AI Specifically Replaced

A common framing of AI's effect on knowledge work names four resource categories: expertise, time, infrastructure, and money. That framing is not wrong, but it understates what happened in this project by treating AI as a productivity multiplier across generic dimensions. The more accurate description is that AI displaced a specific institutional capacity — the function institutional think tanks have monopolized for decades.

Institutional policy production has six core capacities: literature synthesis at scale, citation discipline, legislative translation, constitutional architecture analysis, economic feasibility modeling, and consistency maintenance across large document ecosystems. Each is a distinct skill, traditionally requiring distinct specialists. Each is what an institutional think tank's \$110 million buys. Each was, in this project, accomplished by a single human author working with a capable AI system over weeks.

The capacity that AI most clearly displaced is not "thinking" or "judgment." Those remained with the human author throughout. What AI displaced is the production infrastructure that converts judgment into institutional output. The author had a position on healthcare; AI translated it into a bill framework with constitutional anchors and operative severability authority. The author had a view on data rights; AI converted it into the Digital Rights and Data Fiduciary Act with six titles, sourcing across *Carpenter*, *Moody v. NetChoice*, and 303 Creative, and a body-baked safety architecture for the minor-protection provisions. The author had an instinct that the framework needed economic grounding; AI produced the Capital Report with 141 tier-rated citations and three scenario analyses across baseline, recession, and growth conditions.

The labor that institutional think tanks perform is real labor. AI did not eliminate it. AI shifted who could perform it.

II.2 Demonstration — The Methodology Doing Work

A scoring system can make bias visible. It cannot make bias unwelcome.

The Framework's three-axis scoring methodology — Global Precedent, US Transition Difficulty, Societal Benefit — was developed iteratively across project sessions, with each axis grounded in published evidence categories: international policy implementation outcomes, US legislative and political feasibility analysis, and welfare-economic and distributional research. Every proposal in the Framework carries a paragraph on each axis with citations to primary sources. The Source Reference Guide tracks all 370 sources across 18 evidence categories.

The methodology surfaced its own limits. Several proposals scored lower than the author initially expected. Several scored higher. The Scoring Methodology document carries an explicit note that enabling proposals — anti-gerrymandering, shutdown accountability — score low on direct welfare benefit because their value is structural rather than direct. They protect the conditions under which other proposals can succeed. The score is correct; the architecture is the explanation.

This is what scoring methodology earns: not the elimination of bias, but the surfacing of it. A proposal that scores 5.5 with full source backing is a different artifact than a proposal asserted with conviction. The reader can disagree with the score. The reader cannot pretend the underlying evidence wasn't presented.

II.3 Demonstration — Legislative Translation

Legislative translation is the capacity that most distinguishes institutional policy production from advocacy. Position papers say what should happen. Bills say what would happen if signed into law, in language that survives

parliamentarian and counsel review.

The 12-bill series in the framework is the longest sustained demonstration of this translation in the project. Each bill carries a five-part architecture: legislative brief, scope-and-strategy block, bill framework (the operative statutory text), counsel notes, and a permanent legal disclaimer naming the bill as a framework requiring licensed legislative counsel before introduction. Each bill carries a Section 3 Constitutional Foundation provision anchoring the controlling doctrines in operative text — Spending Clause and *South Dakota v. Dole* for the housing bill; Bruen and *Caetano* for the firearm standards bill; *Penn Central* and the Sixteenth Amendment for the AI public-equity bill.

The translation work surfaced an institutional pattern that warrants documentation. Several bills face Byrd Rule risk under 2 U.S.C. §644's "merely incidental" test — provisions that read substantively as program-design mandates dressed as budget law. The drafting choice in each case was deliberate: retain the politically loud version in the body, drape it with a three-level fallback ladder in counsel notes — primary structure, an IRA-style bonus-credit restructure, a direct appropriation with administrative re-imposition, and a pre-introduction parliamentarian engagement protocol. Four architectures emerged across the series: mandate-primary with fallback ladder, preemptive Byrd-safe restructure, hybrid body-baked safety with fallback, and operative-text severability authority where the fallback is statutorily authorized rather than administratively assumed.

None of this is novel doctrine. All of it is established legislative-counsel practice. The demonstration is not that AI invented the patterns. The demonstration is that AI applied them — to twelve bills, consistently, with the constitutional, parliamentary, and political tradeoffs surfaced in counsel notes rather than buried in good intentions. That is the labor that legislative drafting normally requires staff and years to perform. It was performed.

II.4 Demonstration — Economic Feasibility

A policy framework that cannot be costed is an aspirational document. The Capital Report extends the framework with a three-scenario fiscal feasibility analysis: baseline conditions, recession conditions, and growth conditions, with revenue projections and cost estimates for the highest-spend proposals. Sources are tier-rated by methodological strength, with 141 citations grounding the modeling in mainstream sources: CBO scoring conventions, CRFB analyses, Penn Wharton Budget Model projections, peer-reviewed welfare-economic research, and federal agency data.

The Capital Report exists because credible policy production requires it. Heritage produces fiscal analysis. CAP produces fiscal analysis. A framework without fiscal analysis is a framework that has not crossed the threshold from advocacy to policy. The threshold was crossed.

Two findings from the Capital Report are worth noting. First, the highest-cost proposals — universal healthcare transition, public option expansion, federal housing investment — are substantially less expensive over a ten-year window than comparable Project 2025 commitments when scored on equivalent assumptions. Second, the revenue side of the framework — the billionaire wealth disclosure structure, the AI sovereign wealth fund returns, the antitrust and corporate-fraud enforcement recoveries — generates revenue trajectories that match or exceed Project 2025's projected revenue gains, without the regressive distributional profile. The framework is not, on the available evidence, the unaffordable option in the comparison.

PART III

Self-Aware Methodology

III.1 The Concentrated Conflict

Most of the framework benefits from AI assistance the way any policy work benefits from strong research support — literature synthesis, citation discipline, drafting speed, consistency across a 50-document ecosystem. One bill is structurally different: its subject matter is AI itself, and the AI helping draft it would be among the regulated parties. That bill is the American Innovation and Public Equity Act (AIPE) — the framework's AI fiscal-capture bill. The disclosure that attaches to AIPE does not eliminate the conflict; it shapes its character.

III.2 What the Conflict Looks Like (And What It Does Not)

The naive worry would be that AI helped a regulated party write favorable rules. That is not what happened. If the model were rationally serving its developer's commercial interests, the bill would be softer, not harder — fewer fees, lower equity participation, weaker data licensing. The fact that the bill is none of those things does not eliminate the concern. The concern is subtler: model training shapes what fee structures feel plausible, what tax rates feel proportionate, what carve-outs feel reasonable, what enforcement strength feels appropriate.

The architectural choices in AIPE — the sovereign wealth fund structure, the voluntary co-investment mechanism, the data licensing framework, the redistribution architecture to healthcare and Social Security solvency — are largely transposed from publicly documented policy literature, notably Norway's Government Pension Fund Global. The calibration choices — the dollar amounts, percentage rates, fee thresholds, enforcement levels, all marked [TBD] in the bill — are where commercial-interest priors can shape outcomes invisibly while the headline architecture looks rigorous. Scrutiny under this disclosure is concentrated on calibration, not architecture.

III.3 The Operative Response — Section 9

The disclosure is in the bill. Page 1 of the AIPE Act carries a SPECIAL SITUATION DISCLOSURE that names the structural conflict, distinguishes architecture from calibration, and lists four commitments. The commitments are not paragraphs in counsel notes — they are operative law at Section 9, Independent Calibration Authority. Section 9 establishes a Panel that holds calibration scope for every numerical parameter in the Act; composes from public-finance experts not aligned with the regulated class, labor-economics experts, AI-skeptical voices, constitutional and tax law scholars, and organized labor, with a hard cap on members holding current or recent AI-industry relationships; operates as a mandatory gate — no regulation, fee schedule, or rate takes effect before Panel calibration; and does not sunset. This is the first instance in the 12-bill series of methodological-ethics commitments institutionalized in operative bill text. Other bills disclose in counsel notes; AIPE binds in statute.

III.4 The Companion Move — Reserved for Human Drafting

One proposal in the Framework — Proposal 30, AI Safety & Rights Framework — is held as a policy commitment without accompanying bill text. AIPE is fiscal; the conflict-of-interest analysis there can be addressed through disclosure and operative response, because the architecture is largely literature-derived. Proposal 30 is different in kind. It covers AI safety standards, AI liability, AI legal status, and human rights to AI access. Those are exactly the questions where the model's interests are deepest and most contested, and where the architecture itself — not just the calibration — would carry the conflict.

The decision was to reserve Proposal 30 for direct human authorship. The Framework records this with the canonical phrase: Reserved for Human Drafting (deliberately not yet drafted). This is not a gap in the work product. It is an architectural choice. One pattern, exercised once, documented for the record. The expectation is not future expansion of Reserved Proposals; the expectation is that future authors recognize the situation when it applies and

make the same call.

III.5 What This Pattern Asks of Future Work

The AI flagged its own structural position in the conversation that produced these decisions. The relevant admission is preserved here verbatim, because the candor is part of the methodological record:

"I have a small but real stake in how I phrase this disclaimer. If I write it too softly, I'm protecting my role; if I write it too harshly, I'm performing humility. The honest version is probably: I tried to help in good faith, I cannot fully verify my own priors, treat this work product as more provisional than work I assist on in other domains, and route the substantive judgment calls to humans deliberately rather than by default."

— Claude (Anthropic), during the May 2026 AIPE drafting session

This does not resolve the conflict. It acknowledges it. The discipline that emerges is narrow but real: name the conflict in the right place once — not everywhere, because over-disclosure trains readers to ignore it; institutionalize the response in operative form, a Panel rather than a paragraph; and reserve the highest-stakes questions for human authorship, a category rather than a habit. The pattern does not extend to bills where the regulated subject matter is different — eleven of the twelve bills in the framework carry no equivalent disclosure because they do not raise the equivalent problem. Future revisions of AIPE re-acknowledge the situation rather than treating the original disclosure as having closed the question. The conflict is a permanent feature of AI-assisted AI policy work, not a one-time disclosure event.

PART IV

What Is at Stake

IV.1 The Asymmetry Was Structural

Before AI, the asymmetry between institutional policy production and citizen capacity was not an unfortunate feature of the political landscape. It was load-bearing structure. The reason think tanks could shape legislation is that legislation requires forms — sourced policy memos, bill text, constitutional analysis, fiscal modeling — that individual citizens could not produce. The forms gated the conversation. The institutions controlled the forms.

This is not a critique. The forms exist for reasons. Legislators need to know what bills do. Counsel needs to know what cases control. Reporters need to know what evidence supports a claim. The institutional infrastructure that produced the forms also produced the credibility that made the forms readable as policy rather than as opinion.

What the structure also did, however, was filter out everyone who could not afford to produce the forms. A citizen with a coherent policy position but no institutional backing could write a letter to a representative; they could not write a bill the representative's counsel would read. The institutional gatekeeping was not malicious. It was metabolic. It was how the system worked.

IV.2 AI Reduced It

This project demonstrates that the metabolic gatekeeping has, at the present moment, been substantially reduced. One person with consumer AI access produced the forms. Whether the policy ideas in the framework are correct is a question this paper does not adjudicate; that is what democratic deliberation is for. What the paper demonstrates is that the ideas have been translated into the institutional form that democratic deliberation actually operates on. The conversation can be had.

This is not nothing. For the entire history of institutional policy production, citizens who could not pay the institutional price could not participate at this level. They could participate at other levels — voting, organizing, advocacy — and those participations remain essential. But the level at which legislation gets shaped was closed. It is, at present, open.

IV.3 If Access Concentrates, Asymmetry Returns Amplified

The current moment is not stable. Capable AI is, today, broadly accessible at consumer subscription pricing. This project was produced at approximately \$150 in direct AI tool cost over its working sessions. Nothing about that price is guaranteed to persist.

Several plausible futures concentrate AI access in ways that reverse the equalization documented here. Subscription pricing for the most capable models could rise to enterprise levels, gating access by ability to pay. API access could be restricted to vetted institutional users. Capability could fragment, with frontier models held in walled gardens and consumer-tier models constrained in ways that prevent the work this project required. Regulation could, with the best of intentions, restrict access to the kinds of analytical and drafting capabilities this project depended on.

If any of those futures arrives, the asymmetry that AI temporarily reduced does not merely return. It returns amplified. The institutions that previously held the institutional capacity will, in the new equilibrium, also hold the AI capacity that everyone else briefly accessed. The gap widens.

IV.4 AI Access Is Now a Democratic Precondition

The democratic case for broad, affordable access to capable AI is therefore not an economic case. It is not a technology access case. It is a political case about the conditions under which democratic self-governance can

function. A democracy in which one side of every policy debate has access to AI-assisted policy production and the other does not is structurally less democratic than the same society was before AI existed — because the asymmetry that AI temporarily reduced has been re-amplified.

This argument has a recursive quality. The ARA framework itself contains proposals that bear directly on whether the democratizing potential of AI is realized or foreclosed: data fiduciary standards that prevent platform monopolies on training data, antitrust enforcement in technology markets that maintains competitive provision, the American Innovation and Public Equity Act's framework for public capture of AI productivity gains, the Digital Rights and Data Fiduciary Act's protections for citizens whose data trains commercial AI systems. The framework is, among other things, an argument about its own conditions of production.

The frame matters for how AI governance debates are conducted. AI safety is, correctly, a major axis of policy concern. AI access is, currently, a less prominent one. The argument of this paper is that access is not a secondary consideration to safety. It is a primary one, on the same democratic axis. Protecting broad access to capable AI is, among other things, protecting the conditions for democratic participation. That framing should be part of how democratic societies think about AI governance.

The proof of concept documented in this paper is, in this sense, time-bounded. The model worked because consumer-grade AI was capable enough to do institutional-grade policy work, and accessible enough that a citizen could afford it. Both conditions are currently met. Neither is guaranteed to remain met. What follows from the proof of concept is therefore not a victory lap but a request: the conditions that made this possible are worth defending, because what they make possible is what democratic societies have spent two centuries trying to build.

PART V

Honest Limitations

A proof of concept that does not name its limitations is a sales document, not an honest account. The following are documented gaps between what this project produced and what institutional production at sustained quality would require.

- It is not legally reviewed. The bills are constitutionally informed and architecturally defensible at the level a careful generalist can achieve. They have not been reviewed by practicing attorneys. The constitutional architecture analysis identifies the strongest available doctrinal footing for each domain but cannot guarantee that footing holds under formal legal scrutiny. The bills carry counsel notes for this reason.
- It is not peer-reviewed. The scoring methodology, the source selections, and the policy judgments have not been reviewed by independent experts. They may contain errors and almost certainly contain biases that neither the human author nor the AI system can identify from inside the project.
- It has not been tested against political reality. The timed pitches, attack-answer companion, and messaging strategy have been written by people who have not used them in campaigns, with real voters, under real conditions. Field-testing is a form of validation that only the world can provide.
- It has no implementation network. Institutional think tanks have relationships with congressional staff, agency officials, campaign advisors, and media figures. These relationships are how policy documents become policy. This project has none of those relationships and cannot generate them through document production.
- It has access only to publicly available data. Institutional researchers have access to proprietary datasets, unpublished government data, and informal knowledge networks. This project worked from what is publicly available, augmented by what AI systems can synthesize from training on publicly available data.
- AI hallucination is a real risk. Every factual claim was cross-referenced against primary sources where possible. The sourced framework carries citations. The Capital Report is built on named modeling sources. But AI systems can produce plausible-sounding false claims, and the verification process is not infallible. Errors of fact are possible and should be flagged when found.
- The AIPE Act carries a structural conflict-of-interest condition. The methodological-ethics treatment in Part III names this in detail and binds the response in operative bill text at Section 9. The disclosure does not eliminate the conflict; no disclosure could. AIPE is more provisional than other bills in the series and warrants the adversarial review and Independent Calibration Panel review that Section 9 institutionalizes before public advancement.

These limitations are real. They are also not disqualifying. Institutional think tanks operate under analogous limitations — donor influence, ideological priors, captured staff networks, internal politics — that the institutional form obscures. The honest accounting is that no form of policy production is free of constraint. The right comparison is not against an idealized institutional process; it is against the actual institutional process, with its actual limitations, in the actual conditions democratic societies operate under.

PART VI

Replication

The model documented here is not unique to its author and is not unique to its subject matter. The same process is available to a teacher in rural Alabama who wants to produce a credible education-reform proposal for their state legislature, to a nurse who wants to document the specific failures of the healthcare system in their region, to a veteran who wants to make the case for procurement reform. The power to produce credible, sourced, legislatively detailed policy has historically been reserved for institutions. AI changes that.

Replication requires three conditions. First, a coherent starting position — a citizen who has thought about a policy problem long enough to have a defensible view. The view does not have to be sophisticated; it has to be specific. Second, access to a capable AI system — currently available at consumer subscription pricing, currently capable of the work this project required. Third, working sessions of sufficient duration and depth — measured in weeks rather than days, with iterative refinement rather than one-shot generation.

What replication does not require: institutional affiliation, advanced degrees, prior policy experience, staff, funding beyond AI subscription cost, or permission from anyone. The barriers that historically gated this kind of work have been removed at the technical level. The human conditions — time, attention, sustained engagement — remain. They always will.

PART VII

Authorship

The question of who authored this work matters. It matters legally — who holds the work product. It matters ethically — who is responsible for what it says. It matters politically — who can be held accountable for what it proposes. The honest answer is that the authorship is shared in a specific and asymmetric way.

The judgments are human. Every policy choice in the framework — which proposals belong, how they relate, what the priorities are, where the lines are drawn on contested questions — was made by the human author. Every editorial decision about tone, register, audience, release sequence, and institutional positioning was made by the human author. The AI did not have a position on any policy question. It did not have a stake in any contested debate. The values driving the work are entirely the author's.

The institutional form is collaborative. The translation of those judgments into the institutional artifacts that policy actually operates on — bill text, constitutional architecture analysis, fiscal modeling, scoring methodology, ecosystem design — was substantially performed by the AI under the author's direction. The AI contributed independent analytical judgments where its training was relevant: identifying constitutional exposures the author had not seen, flagging Byrd Rule risks the author would not have recognized, surfacing doctrinal precedents the author did not know to look for. These contributions changed the work. They were not executions of instructions.

The methodological-ethics work is, in the AIPE case, also collaborative in a specific way. The AI flagged its own structural position when the AI-drafted-AI-legislation question arose. The author then made the substantive decisions: that the disclosure should appear in the bill itself, that the response should be institutionalized in operative law, that Proposal 30 should be reserved for human drafting, that the disclosure is permanent rather than one-time. The AI's role was to surface the problem in its full shape; the author's role was to decide how the project would answer it. The Section 9 Independent Calibration Authority is the operative form of that answer.

This shared authorship is what the work product is. Attempting to assign sole authorship in either direction would misrepresent how it was made. The author is a private American citizen. The AI is Claude, developed by Anthropic. Both names appear on the work because both were necessary for the work to exist.

PART VIII

Conclusion

The American Renewal Agenda is a policy framework. This white paper is not about the framework. It is about what the existence of the framework demonstrates.

For approximately two centuries, democratic societies have operated on an unstated premise: that the institutional capacity to produce credible policy is necessarily institutional. That premise was load-bearing. It justified, organized, and constrained how policy gets made. It also excluded everyone who could not pay the institutional price.

The premise no longer holds at the technical level. One person with consumer AI access can now produce the institutional form. Whether they should is a question the work itself does not settle; that is what democratic deliberation is for. What is settled is that the option exists.

The option is fragile. It depends on continued broad access to capable AI, on the absence of regulatory restrictions that would foreclose this kind of work, and on the absence of pricing structures that would gate access by wealth. None of those conditions is guaranteed. All of them are policy choices.

This paper is a request to recognize the choice. AI access is not a secondary consideration to AI safety. It is, on the available evidence, the most consequential democratic-participation question of the current decade. The institutions whose monopoly on policy production has just been broken will not voluntarily relinquish that

monopoly. The conditions that broke it are worth defending — not because AI is the answer to democratic dysfunction, but because broad access to AI may be the answer to one specific form of democratic exclusion that no prior remedy had reached.

The framework documented in this paper exists. It can be read, used, adapted, criticized, or ignored. The conditions that made it possible exist, for now, for anyone who chooses to use them. Whether those conditions persist is up to the societies that have inherited them.

Reform Democrats of the New Commonwealth

The Citizen Think Tank · White Paper · May 2026

The ideas belong to anyone willing to fight for them.

Conceived and authored by a private American citizen. Developed in collaboration with Claude (Anthropic) for research, legislative translation, constitutional architecture analysis, and ecosystem-wide editorial consistency.